

Why I rewrote my capstone project in BSV

Author: Sai Govardhan M C

Email: sai.govardhan@incoresemi.com

Introduction

This document outlines the motivation for using Bluespec System Verilog, which provides high levels of abstraction to rapidly design hardware microarchitecture.

As with most ECE/CSE students, I started out designing initial digital logic in Verilog, with reference to coding guidelines from professors, coursework, and the [Sunburst Papers](#).

Specifying designs with parallelism and concurrency in Verilog was always a challenge due to lower levels of abstraction and regular rework to fix synthesis-simulation mismatches.

At InCore, the use of BSV is one of the superpowers that enables small teams like ours to specify complex hardware intuitively, correctly, and efficiently. A year ago, as a novice BSV designer, I decided it would be meaningful to contrast the efforts that my team and I spent during our capstone project by re-implementing the Multi-Dimensional Sorting Algorithm (MDSA) in BSV. This blog collates these insights to establish a strong use-case for teaching BSV at universities.

In this post, I shall be diving into the BSV implementation by explaining essential parts of the microarchitecture, the ease of specifying them in BSV, and corresponding snippets from the codebase.

More of our work on the taxonomy of sorters, low power methodologies, other variants (Hybrid and Odd-Even sorters), and our ASIC implementation results can be referred to in our published paper [Low Power Multidimensional Sorters using Clock Gating and Index Sorting](#).

My complete MDSA Bitonic Implementation in BSV, along with our legacy Verilog implementation, can be found in my [GitHub repository](#).

The Compare And Exchange Block

The **Compare And Exchange** (CAE) block is a fundamental building block of Systolic Array based Parallel Hardware Sorters which sorts two inputs to an ascending order output.




Figure 1. The Compare and Exchange (CAE) Block

- Specify the CAE typedef as a Vector of 2 elements:

- Declare the method ActionValue `mav_get_sort`:

The CAE block checks if `cae_in[0]` is greater than `cae_in[1]` and uses the inbuilt Vector to Vector `reverse` function to swap the values.

```
method ActionValue#(CAE) mav_get_sort (CAE cae_in);
    if(cae_in[0] > cae_in[1]) begin
        cae_in = reverse(cae_in);
    end
    return(cae_in);
endmethod
```

The Bitonic Sorting Unit

The Bitonic Sorting Unit is a network of 24 such CAE blocks, arranged as depicted below. This network sorts eight input elements in ascending order at the end of six stages.




Figure 2. The Bitonic Merge Sorting Network

The BM4 sorter

We could easily modularize a part of this design, the BM4 unit, by creating an intermediate two-stage, four-input sorter, and specify the two methods for input and output as follows:




Figure 3. The BM4 sorter Network

- Declare a typedef for BM4 as a vector of 4 inputs:

- Specify the intermediate pipeline as a register of the BM4 type:

```
Reg#(BM4) pipe <- mkReg(unpack(0));
```

- First stage of sorting with the inputs, by routing the inputs at indices 0 and 3 to CAE-0, and 1 and 2 to CAE-1 block. We use the inbuilt vector function `vec()` to combine multiple elements into a vector:

```
let lv_get_sort_1 <- cae[0].mav_get_sort(vec(bm4[0], bm4[3]));
let lv_get_sort_2 <- cae[1].mav_get_sort(vec(bm4[1], bm4[2]));
// Store intermediate results in a pipeline
pipe <= vec(lv_get_sort_1[0], lv_get_sort_2[0], lv_get_sort_2[1],
lv_get_sort_1[1]);
```

- Perform the second stage sorting with the intermediate sorted values by routing the pipeline outputs at indices 0 and 1 to CAE-0, and 2 and 3 to CAE-1 block:

```
let lv_get_sort_3 <- cae[0].mav_get_sort(vec(pipe[0], pipe[1]));
let lv_get_sort_4 <- cae[1].mav_get_sort(vec(pipe[2], pipe[3]));
```

- Return the outputs as:

```
return (vec(lv_get_sort_3[0], lv_get_sort_3[1], lv_get_sort_4[0],
lv_get_sort_4[1]));
```

The BM8 sorter

Now with the abstraction of using a BM4 sorter, we can proceed to design the complete Bitonic Merge 8 input sorter as follows:

- Instantiate the 5 intermediate register pipelines:

```
Vector#(5, Reg#(BM8)) pipe <- replicateM(mkReg(unpack(0)));
```

- Pass the inputs through the network defined for each stage of the BM8, while storing the intermediate values in the above pipeline registers:
- Stage 1:

```

let lv_cae_sort_1 <- cae_stage_1[0].mav_get_sort(vec(bm8_in[0], bm8_in[1]));
let lv_cae_sort_2 <- cae_stage_1[1].mav_get_sort(vec(bm8_in[2], bm8_in[3]));
let lv_cae_sort_3 <- cae_stage_1[2].mav_get_sort(vec(bm8_in[4], bm8_in[5]));
let lv_cae_sort_4 <- cae_stage_1[3].mav_get_sort(vec(bm8_in[6], bm8_in[7]));

```

```

pipe[0] <= vec(lv_cae_sort_1[0]
               , lv_cae_sort_1[1]
               , lv_cae_sort_2[0]
               , lv_cae_sort_2[1]
               , lv_cae_sort_3[0]
               , lv_cae_sort_3[1]
               , lv_cae_sort_4[0]
               , lv_cae_sort_4[1]);

```

- Stage 2:

Pass the outputs of the first stage to the BM4 sorter, and register their output for the third stage:

```

bm4_stage_2_3[0].ma_get_inputs(vec(pipe[0][0], pipe[0][1], pipe[0][2],
pipe[0][3]));
bm4_stage_2_3[1].ma_get_inputs(vec(pipe[0][4], pipe[0][5], pipe[0][6],
pipe[0][7]));

```

```

pipe[1] <= vec(lv_get_bm4_sort_1[0]
               , lv_get_bm4_sort_1[1]
               , lv_get_bm4_sort_1[2]
               , lv_get_bm4_sort_1[3]
               , lv_get_bm4_sort_2[0]
               , lv_get_bm4_sort_2[1]
               , lv_get_bm4_sort_2[2]
               , lv_get_bm4_sort_2[3]);

```

... and so on for the remaining stages.

The MDSA Algorithm Implementation

The MDSA algorithm efficiently uses Parallel Hardware Sorters (PHSAs) like the Bitonic sorter we earlier designed to specify an architecture that uses eight such units to sort 64 elements in 6 stages by only alternating between row and column sorting, and rerouting the order of outputs (ascending/descending).

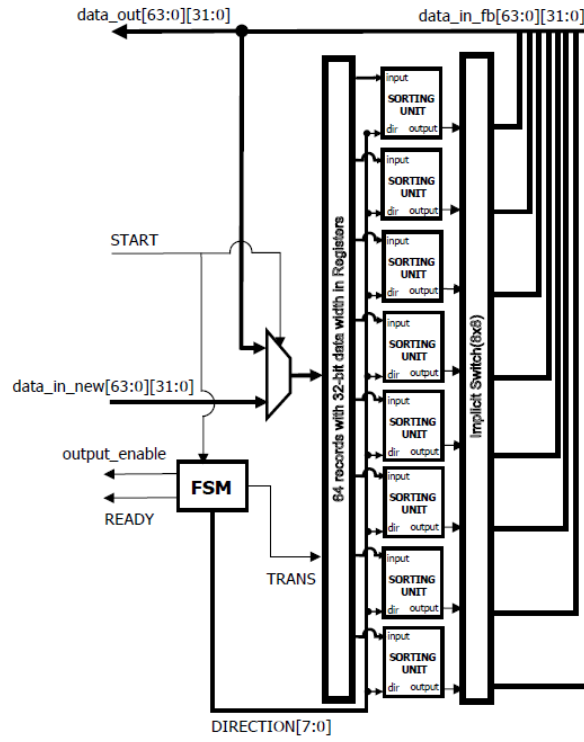


Figure 4. The block diagram of the MDSA Architecture

MDSA Algorithm FSM

Phase	Row/ Column Sorting	Ascending sorting	Descending sorting
1	Column	1,2,3,4,5,6,7,8	-
2	Row	1,3,5,7	2,4,6,8
3	Column	1,2,3,4,5,6,7,8	-
4	Row	2,4,6,8	1,3,5,7
5	Column	1,2,3,4,5,6,7,8	-
6	Row	1,2,3,4,5,6,7,8	-

Figure 5. The FSM that implements the MDS-Algorithm

We specify the **MDSA_64** type which is a multidimensional 8x8 vector

```
typedef Vector#(8, Vector#(8, Bit#(WordLength))) MDSA_64;
```

To create a 64 record register buffer specified as:

```
Reg#(MDSA_64) v_rg_mdsa_in <- mkReg(unpack(0));
```

And use this helper function to send inputs to the MDSA sorter network:

```

function fn_input_sorting_network(Vector#(8, Ifc_bm8) bm8
                                , MDSA_64 mdsa_in);

    action
        bm8[0].ma_get_inputs(mdsa_in[0]);
        bm8[1].ma_get_inputs(mdsa_in[1]);
        bm8[2].ma_get_inputs(mdsa_in[2]);
        bm8[3].ma_get_inputs(mdsa_in[3]);
        bm8[4].ma_get_inputs(mdsa_in[4]);
        bm8[5].ma_get_inputs(mdsa_in[5]);
        bm8[6].ma_get_inputs(mdsa_in[6]);
        bm8[7].ma_get_inputs(mdsa_in[7]);
    endaction
endfunction

```

Stage 1: Column Sorting

- Sending the inputs to the Eight BM8 sorters:

```

rule rl_mdsa_send_inputs_to_stage_1(rg_mdsa_fsm == STAGE_1_IN);
    // Column Sorting Phase
    $display("[MDSA] STARTING MDSA STAGE 1");
    $display("[MDSA]: STAGE 1 INPUTS:", fshow(v_rg_mdsa_in));
    fn_input_sorting_network(bm8, v_rg_mdsa_in);
    rg_mdsa_fsm <= STAGE_1_OUT;
endrule

```

- Collecting the ascending order of responses

```

lv_s1_output[0] <- bm8[0].mav_return_outputs();
lv_s1_output[1] <- bm8[1].mav_return_outputs();
lv_s1_output[2] <- bm8[2].mav_return_outputs();
lv_s1_output[3] <- bm8[3].mav_return_outputs();
lv_s1_output[4] <- bm8[4].mav_return_outputs();
lv_s1_output[5] <- bm8[5].mav_return_outputs();
lv_s1_output[6] <- bm8[6].mav_return_outputs();
lv_s1_output[7] <- bm8[7].mav_return_outputs();

```

- Transposing the output: We can use the inbuilt transpose function in BSV to alternate between the row and column sorting between the phases of the MDSA.

```

v_rg_mdsa_in <= transpose(lv_s1_output);

```

Stage 2: Row Sorting

- Sending the inputs to the Eight BM8 sorters:

```

rule rl_mdsa_send_inputs_to_stage_2(rg_mdsa_fsm == STAGE_2_IN);
  // Row Sorting Phase
  $display("[MDSA] STARTING MDSA STAGE 2");
  $display("[MDSA]: STAGE 2 INPUTS:", fshow(v_rg_mdsa_in));
  fn_input_sorting_network(bm8, v_rg_mdsa_in);
  rg_mdsa_fsm <= STAGE_2_OUT;
endrule

```

- Collecting the alternating ascending and descending order of responses

```

lv_s2_output[0] <- bm8[0].mav_return_outputs();
lv_s2_output[1] <- bm8[1].mav_return_outputs();
lv_s2_output[1] = reverse(lv_s2_output[1]);
lv_s2_output[2] <- bm8[2].mav_return_outputs();
lv_s2_output[3] <- bm8[3].mav_return_outputs();
lv_s2_output[3] = reverse(lv_s2_output[3]);
lv_s2_output[4] <- bm8[4].mav_return_outputs();
lv_s2_output[5] <- bm8[5].mav_return_outputs();
lv_s2_output[5] = reverse(lv_s2_output[5]);
lv_s2_output[6] <- bm8[6].mav_return_outputs();
lv_s2_output[7] <- bm8[7].mav_return_outputs();
lv_s2_output[7] = reverse(lv_s2_output[7]);

```

Transposing the output

```

v_rg_mdsa_in <= transpose(lv_s2_output);

```

... and so on for the remaining stages as per the MDSA Algorithm FSM.

Ultimately, an ideal test case where all 64 inputs specified in descending order:

```

[MDSA] STARTING MDSA STAGE 1
[MDSA]: STAGE 1 INPUTS:<V <V 'h00000040 'h0000003f 'h0000003e 'h0000003d 'h0000003c
'h0000003b 'h0000003a 'h00000039 > <V 'h00000038 'h00000037 'h00000036 'h00000035
'h00000034 'h00000033 'h00000032 'h00000031 > <V 'h00000030 'h0000002f 'h0000002e
'h0000002d 'h0000002c 'h0000002b 'h0000002a 'h00000029 > <V 'h00000028 'h00000027
'h00000026 'h00000025 'h00000024 'h00000023 'h00000022 'h00000021 > <V 'h00000020
'h0000001f 'h0000001e 'h0000001d 'h0000001c 'h0000001b 'h0000001a 'h00000019 > <V
'h00000018 'h00000017 'h00000016 'h00000015 'h00000014 'h00000013 'h00000012
'h00000011 > <V 'h00000010 'h0000000f 'h0000000e 'h0000000d 'h0000000c 'h0000000b
'h0000000a 'h00000009 > <V 'h00000008 'h00000007 'h00000006 'h00000005 'h00000004
'h00000003 'h00000002 'h00000001 > >

```

Shall be sorted in 6 stages to ascending order as follows:

```
Final MDSA output: <%h><V <V 'h00000001 'h00000002 'h00000003 'h00000004 'h00000009
'h0000000a 'h0000000b 'h0000000c > <V 'h00000005 'h00000006 'h00000007 'h00000008
'h0000000d 'h0000000e 'h0000000f 'h00000010 > <V 'h00000011 'h00000012 'h00000013
'h00000014 'h00000019 'h0000001a 'h0000001b 'h0000001c > <V 'h00000015 'h00000016
'h00000017 'h00000018 'h0000001d 'h0000001e 'h0000001f 'h00000020 > <V 'h00000021
'h00000022 'h00000023 'h00000024 'h00000029 'h0000002a 'h0000002b 'h0000002c > <V
'h00000025 'h00000026 'h00000027 'h00000028 'h0000002d 'h0000002e 'h0000002f
'h00000030 > <V 'h00000031 'h00000032 'h00000033 'h00000034 'h00000039 'h0000003a
'h0000003b 'h0000003c > <V 'h00000035 'h00000036 'h00000037 'h00000038 'h0000003d
'h0000003e 'h0000003f 'h00000040 > >
Verilog simulation finished
```

References

Steps to run simulations of the CAE, BM4, BM8 and MDSA_Bitonic in the GitHub Repository

Acknowledgements